

# MINI PROJECT

## PROBLEM STATEMENT: which model is suitable for Insurance Dataset

### Importing Packages

```
In [1]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

### Read the data

```
In [2]: df=pd.read_csv(r"C:\Users\arshiha\Downloads\insurance.csv")
df
```

```
Out[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

## Data Collection And Preprocessing

```
In [3]: df.head()
```

```
Out[3]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [4]: df.tail()
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1338 entries, 0 to 1337  
Data columns (total 7 columns):  
#   Column      Non-Null Count  Dtype    
---  ---        
0   age         1338 non-null   int64    
1   sex         1338 non-null   object   
2   bmi         1338 non-null   float64  
3   children    1338 non-null   int64    
4   smoker      1338 non-null   object   
5   region      1338 non-null   object   
6   charges     1338 non-null   float64  
dtypes: float64(2), int64(2), object(3)  
memory usage: 73.3+ KB
```

```
In [6]: df.shape
```

```
Out[6]: (1338, 7)
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [8]: df.isna().any()
```

```
Out[8]: age      False
sex        False
bmi        False
children   False
smoker     False
region     False
charges    False
dtype: bool
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: age      0
sex        0
bmi        0
children    0
smoker     0
region     0
charges    0
dtype: int64
```

```
In [10]: df.fillna(method="ffill",inplace=True)
```

```
In [11]: x=np.array(df["age"]).reshape(-1,1)
```

```
In [12]: y=np.array(df["children"]).reshape(-1,1)
```

```
In [13]: df.dropna(inplace=True)
```


## Linear Regression

```
In [14]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```

```
In [15]: from sklearn.linear_model import LinearRegression  
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```

-0.0017866536487798346

```
In [16]: from sklearn import preprocessing, svm  
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color="r")  
plt.plot(x_test,y_pred,color="k")  
plt.show()
```

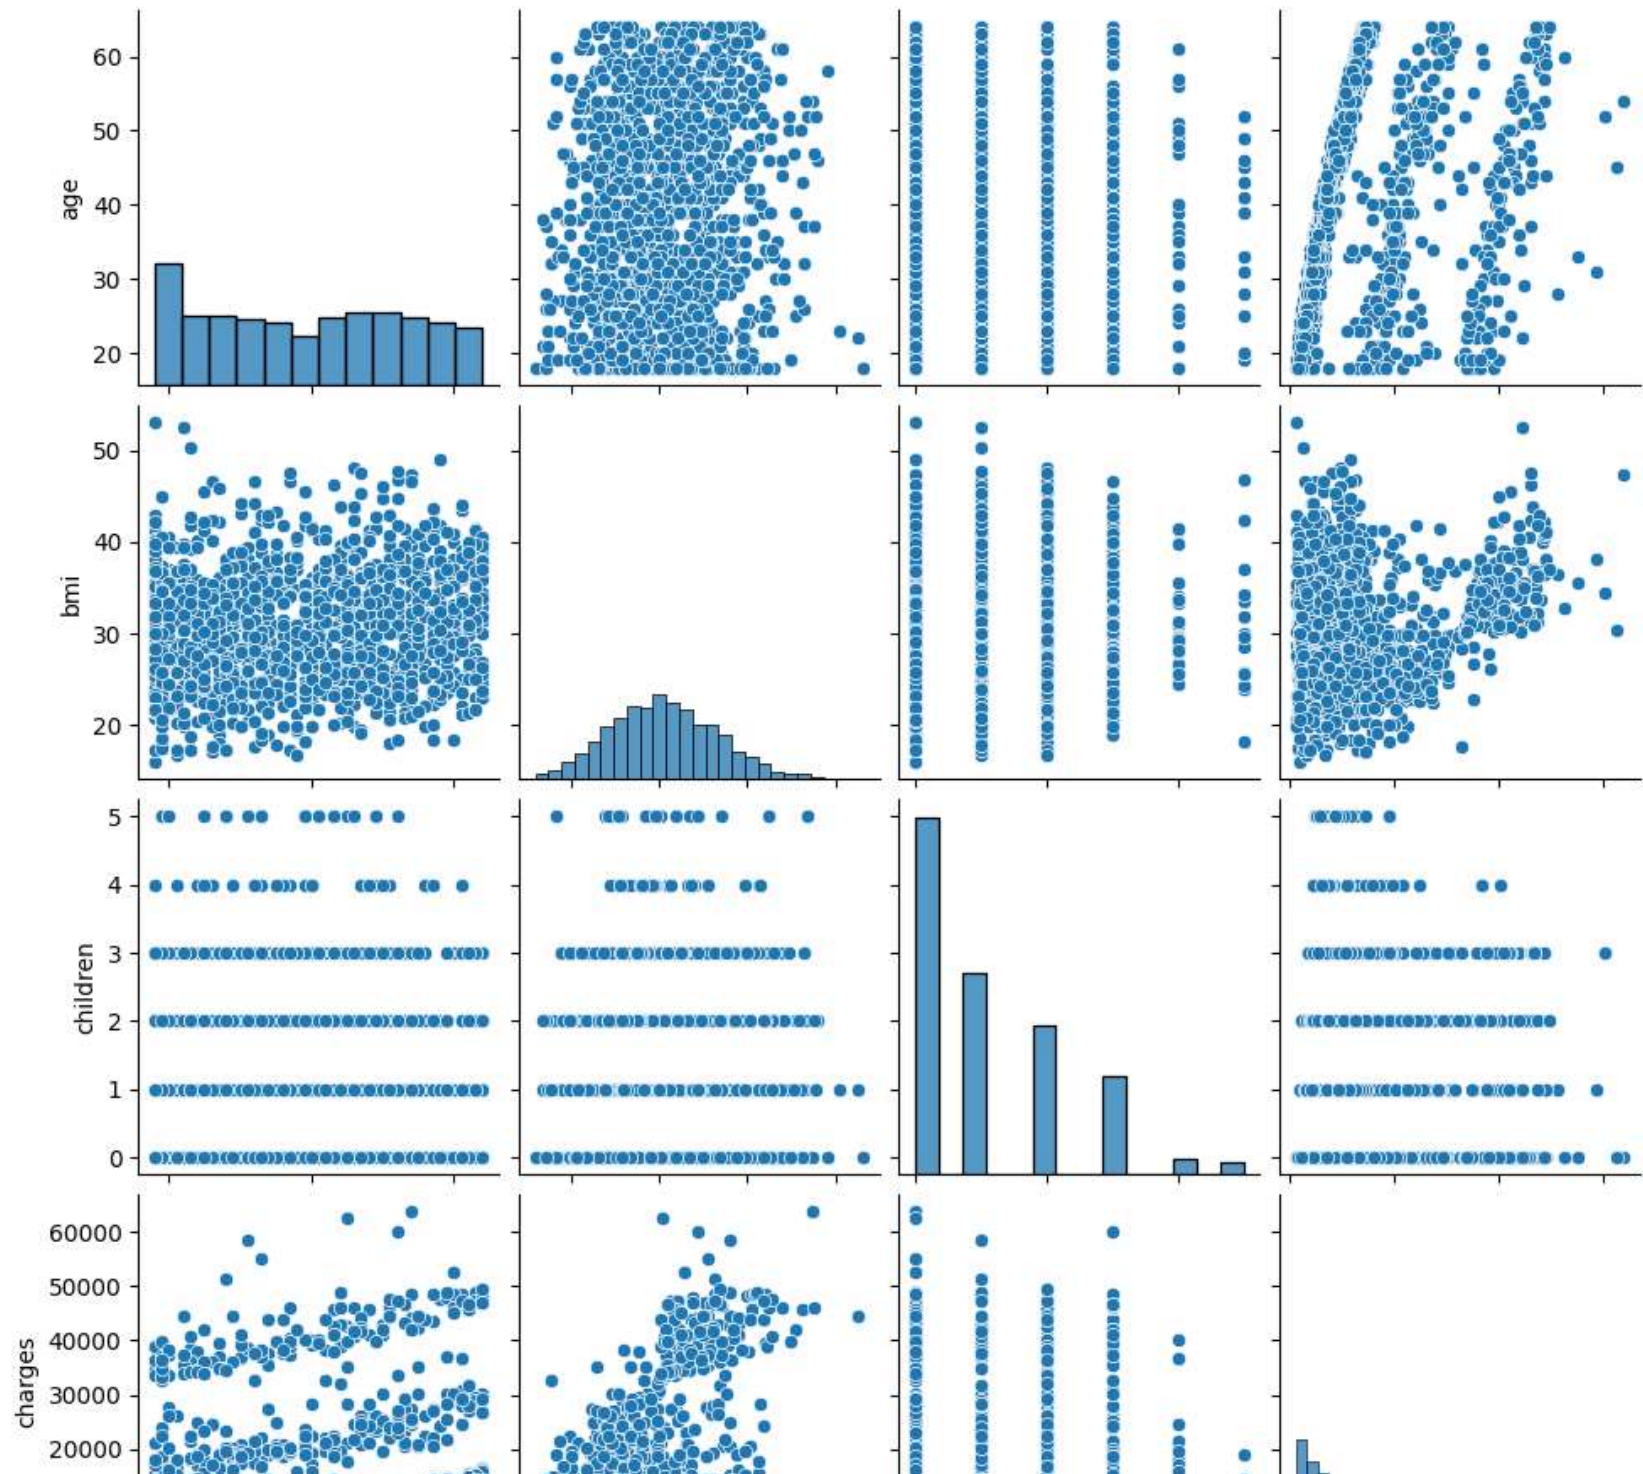



```
In [17]: sns.pairplot(df)
```

```
Out[17]: <seaborn.axisgrid.PairGrid at 0x212fab326d0>
```











```
In [18]: sns.displot(df['age'])
```

```
Out[18]: <seaborn.axisgrid.FacetGrid at 0x212ffc64250>
```



```
In [19]: sns.displot(df['children'])
```

```
Out[19]: <seaborn.axisgrid.FacetGrid at 0x212ff639b50>
```




```
In [20]: plt.figure(figsize=(15,8))
```

```
Out[20]: <Figure size 1500x800 with 0 Axes>
```

```
<Figure size 1500x800 with 0 Axes>
```


```
In [21]: df500=df[:][:500]  
sns.lmplot(x="age",y="children",data=df500,order=1,ci=None)
```

```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x212ffc380d0>
```



```
In [22]: df500.fillna(method='ffill',inplace=True)
x=np.array(df500['age']).reshape(-1,1)
y=np.array(df500['children']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: -0.0042836432283639425



```
In [23]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("r2 score:",r2)
```

r2 score: -0.0042836432283639425

## Decision Tree

```
In [24]: from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

```
In [25]: from sklearn.tree import DecisionTreeClassifier
```

```
In [26]: df['sex'].value_counts()
```

```
Out[26]: sex
male      676
female    662
Name: count, dtype: int64
```

```
In [27]: df['smoker'].value_counts()
```

```
Out[27]: smoker
no      1064
yes      274
Name: count, dtype: int64
```

```
In [28]: df['region'].value_counts()
```

```
Out[28]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

```
In [29]: convert={"sex":{"female":1,"male":2}}
df=df.replace(convert)
df
```

```
Out[29]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	2	33.770	1	no	southeast	1725.55230
2	28	2	33.000	3	no	southeast	4449.46200
3	33	2	22.705	0	no	northwest	21984.47061
4	32	2	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	2	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500
1337	61	1	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [30]: convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

```
Out[30]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	2	33.770	1	0	southeast	1725.55230
2	28	2	33.000	3	0	southeast	4449.46200
3	33	2	22.705	0	0	northwest	21984.47061
4	32	2	28.880	0	0	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	2	30.970	3	0	northwest	10600.54830
1334	18	1	31.920	0	0	northeast	2205.98080
1335	18	1	36.850	0	0	southeast	1629.83350
1336	21	1	25.800	0	0	southwest	2007.94500
1337	61	1	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns



```
In [31]: convert={"region":{"southwest":1,"southeast":2,"northwest":3,"northeast":4}}
df=df.replace(convert)
df
```

```
Out[31]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	1	16884.92400
1	18	2	33.770	1	0	2	1725.55230
2	28	2	33.000	3	0	2	4449.46200
3	33	2	22.705	0	0	3	21984.47061
4	32	2	28.880	0	0	3	3866.85520
...	...	...	...	...	...	...	...
1333	50	2	30.970	3	0	3	10600.54830
1334	18	1	31.920	0	0	4	2205.98080
1335	18	1	36.850	0	0	2	1629.83350
1336	21	1	25.800	0	0	1	2007.94500
1337	61	1	29.070	0	1	3	29141.36030

1338 rows × 7 columns

```
In [32]: x=["age","bmi","children"]
y=[1,2]
all_inputs=df[x]
all_classes=df["sex"]
```

```
In [33]: (x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.25)
```

```
In [34]: clf=DecisionTreeClassifier(random_state=0)
```

```
In [35]: clf.fit(x_train,y_train)
```

```
Out[35]: DecisionTreeClassifier(random_state=0)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [36]: `clf.score(x_test,y_test)`

Out[36]: 0.4835820895522388


## Data Visualization

```
In [37]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

```
Out[37]: <Axes: >
```









```
In [38]: df.drop(columns=["region", "bmi"], inplace=True)
sns.pairplot(df)
df.smoker=np.log(df.smoker)
```

```
C:\Users\arshiha\AppData\Local\Programs\Python\Python311\Lib\site-packages\pandas\core\arraylike.py:396: RuntimeWarning: divide by zero encountered in log
  result = getattr(ufunc, method)(*inputs, **kwargs)
```









```
In [39]: features=df.columns[0:1]
target=df.columns[-1]
x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
Scaler=StandardScaler()
x_train=Scaler.fit_transform(x_train)
x_test=Scaler.transform(x_test)
```

The dimension of x\_train is (936, 1)  
The dimension of x\_test is (402, 1)

```
In [40]: lr = LinearRegression()
lr.fit(x_train,y_train)
actual = y_test
train_score_lr = lr.score(x_train,y_train)
test_score_lr = lr.score(x_test,y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07447061146193878  
The test score for lr model is 0.10891203216512224

```
In [41]: ridgeReg = Ridge(alpha=10)
         ridgeReg.fit(x_train,y_train)
         train_score_ridge = ridgeReg.score(x_train, y_train)
         test_score_ridge = ridgeReg.score(x_test, y_test)
         print("\nRidge Model:\n")
         print("The train score for ridge model is {}".format(train_score_ridge))
         print("The test score for ridge model is {}".format(test_score_ridge))
```


Ridge Model:

The train score for ridge model is 0.07446228994221393

The test score for ridge model is 0.10855133360950642

```
In [42]: plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='LinearReg')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```





age

## Lasso Regression

```
In [43]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.07446997086306062
0.10881427793326703
```


```
In [44]: print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(x_train,y_train)
train_score_ls =lasso.score(x_train,y_train)
test_score_ls =lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

```
The train score for ls model is 0.07446997086306062
The test score for ls model is 0.10881427793326703
```

```
In [45]: pd.Series(lasso.coef_,features).sort_values(ascending = True).plot(kind = "bar")
```

```
Out[45]: <Axes: >
```



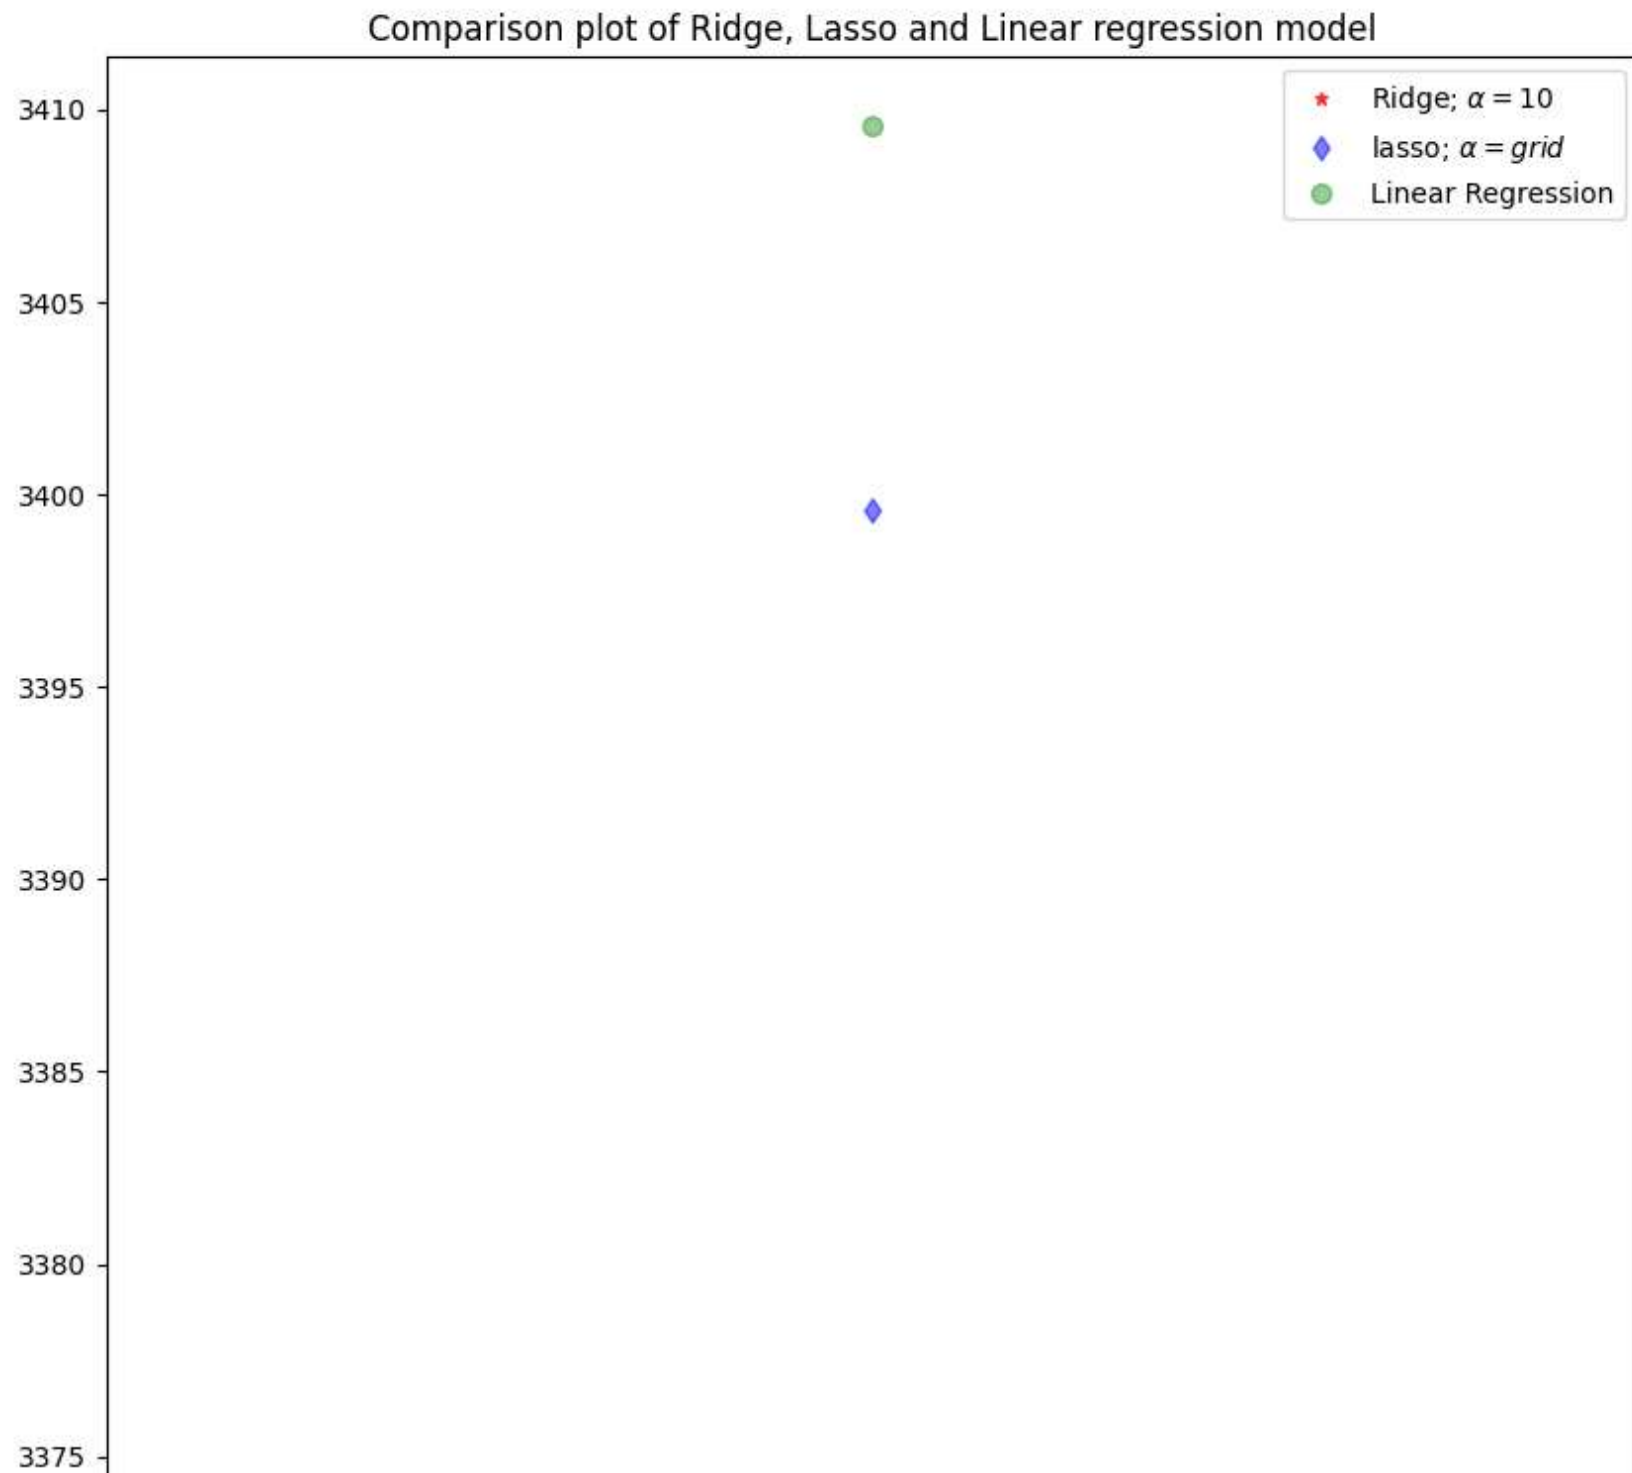
```
In [46]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```


```
0.07446997086306062
0.10881427793326703
```



```
In [47]: plt.figure(figsize=(10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\alpha$ ')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Reg')
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```







```
In [48]: from sklearn.linear_model import RidgeCV
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(x_train, y_train)
print("The train score for ridge model is {}".format(ridge_cv.score(x_train, y_train)))
print("The test score for ridge model is {}".format(ridge_cv.score(x_test, y_test)))
```

The train score for ridge model is 0.07446228994221393

The test score for ridge model is 0.10855133360950775

## Elastic Net

```
In [49]: from sklearn.linear_model import ElasticNet
regr = ElasticNet()
regr.fit(x, y)
print(regr.coef_)
print(regr.intercept_)
```

[257.0684655]  
3191.532406056682

```
In [50]: y_pred_elastic = regr.predict(x_train)
```

```
In [51]: mean_squared_error = np.mean((y_pred_elastic - y_train)**2)
print("Mean Squared Error on test set", mean_squared_error)
```

Mean Squared Error on test set 267460995.25217086

```
In [52]: import re
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
%matplotlib inline
digits=load_digits()
```

```
In [53]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_size=0.7,random_state=42)
```

```
In [54]: from sklearn.linear_model import LogisticRegression
```

```
In [55]: logisticRegr=LogisticRegression(max_iter=10000)
logisticRegr.fit(x_train,y_train)
```

Out[55]: LogisticRegression(max\_iter=10000)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [56]: print(logisticRegr.predict(x_test))

[6 9 3 ... 1 0 8]
```

```
In [57]: score=logisticRegr.score(x_test,y_test)
print(score)

0.9467408585055644
```

## Random Forest

```
In [58]: from sklearn.ensemble import RandomForestClassifier  
rf=RandomForestClassifier()  
rf.fit(x_train,y_train)
```

Out[58]: RandomForestClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [59]: rf=RandomForestClassifier()
```

```
In [60]: params={'max_depth':[2,3,4,5,6],  
               'min_samples_leaf':[5,10,15,20,50,100],  
               'n_estimators':[10,25,30,50,100,200]}
```

```
In [ ]: from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
In [66]: grid_search.best_score_
```

Out[66]: 0.9054041029877461

```
In [67]: rf_best=grid_search.best_estimator_  
print(rf_best)
```

RandomForestClassifier(max\_depth=6, min\_samples\_leaf=5)

```
In [68]: x=df.drop("smoker",axis=1)  
y=df["smoker"]
```

## CONCLUSION

**Based on the accuracy scores of all models we can conclude that "Logistic Regression" is the best model for the given data set.**

